

RPC: Remote Procedure Call

February 23, 2018

Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Roadmap

Idea

Implementation

Transparency


RPC Semantics in the Presence of Faults

Remote Procedure Call (RPC)


- ▶ Message-based programming with `send()`/`receive()` primitives is not convenient
 - ▶ depends on the communication protocol used (TCP vs. UDP)
 - ▶ requires the specification of an application protocol
 - ▶ akin to I/O
- ▶ Function/procedure call in a remote computer
 - ▶ is a familiar paradigm
 - ▶ eases transparency
 - ▶ is particularly suited for client-server applications

RPC: the Idea

Local procedure call:




Remote procedure call:




Program Development with RPCs: the Vision

- Design/develop an application ignoring distribution



- Distribute *a posteriori*



Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults


RPC Stub Routines

- ▶ Ensure RPC transparency


Client invokes the **client stub** – a local function

Remote function is invoked by the **server stub** – a local function


- ▶ The stub routines communicate with one another by exchanging messages



Well Known Trick: also Used for System Calls



Typical Architecture of an RPC System



Obs. RPC is typically implemented on top of the transport layer (TCP/IP)

Client Stub

Request

1. Assembles message: **parameter marshalling**
2. Sends message, via `write()` / `sendto()` to server
3. Blocks waiting for response, via `read()` / `recvfrom()`
 - ▶ Not in the case of **asynchronous RPC**

Response

1. Receives responses
2. Extracts the results (**unmarshalling**)
3. Returns to client
 - ▶ Assuming **synchronous RPC**

Server Stub

Request


1. Receives message with request, via `read()` / `recvfrom()`
2. Parses message to determine arguments (**unmarshalling**)
3. Calls function

Response

1. Assembles message with the return value of the function
2. Sends message, via `write()` / `sendto()`
3. Blocks waiting for a new request

RPC: Dispatching

- ▶ Often, **RPC services** offer more than one remote procedure:



- ▶ The identification of the procedure is performed by the **dispatcher**
 - ▶ This leads to a hierarchical name space (**service, procedure**)

Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Transparency: Platform Heterogeneity

Problems at least two:

1. Different architectures use different formats
 - ▶ 1's-complement vs. 2's complement
 - ▶ big-endian vs. little-endian
 - ▶ ASCII vs. UTF-??
2. Languages or compilers may use different representations for composite data-structures

Solution mainly two:

standardize format in the wires

- + needs only two conversions in each platform
- may not be efficient

receiver-makes-right

Transparency: Addresses as Arguments

Issue The meaning of an address (C pointer) is specific to a process

Solution Use **call-by-copy/restore** for parameter passing

- + Works in most cases
- Complex
 - ▶ The same address may be passed in different arguments
- Inefficient
 - ▶ For complex data structures, e.g. trees

Transparency in the Presence of Faults

Problem What if something breaks?

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
 - ▶ Was the request processed before the crash?

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
 - ▶ Was the request processed before the crash?
- ▶ Client crashes

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
 - ▶ Was the request processed before the crash?
- ▶ Client crashes
 - ▶ Need to prevent **orphan** computations, i.e. on behalf of a dead process.

Transparency in the Presence of Faults

Problem What if something breaks?

- ▶ The client cannot locate the server
 - ▶ RPC can return an error (like in the case of a system call)
- ▶ The request-message is lost
 - ▶ Retransmit it, after a timeout
- ▶ The response-message is lost
 - ▶ Must use request identifiers (sequence nos.)
 - ▶ Must save most recent responses for replay, if the request is not **idempotent**
- ▶ Server crashes
 - ▶ Was the request processed before the crash?
- ▶ Client crashes
 - ▶ Need to prevent **orphan** computations, i.e. on behalf of a dead process.

Issue A client cannot distinguish between loss of a request, loss of a response or a server crash

- ▶ The absence of a response may be caused by a slow network/server

Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

RPC Semantics in the Presence of Faults (Spector82)

Question What can a client expect when there is a fault?

Answer Depends on the semantics in the presence of faults provided by the RPC system

At-least-once Client stub must keep retransmitting until it obtains a response


- ▶ Be careful with non-idempotent operations
- ▶ Spector allows for zero executions in case of server failure

At-most-once Not trivial if you use a non-reliable transport, e.g. UDP.


- ▶ If the RPC uses TCP, it may report an error when the TCP connection breaks

Exactly-once Not always possible to ensure this semantics, especially if there are external actions that cannot be undone


Faults and Exactly-once Semantics



(a)



(b)



(c)

Problem In the case of external actions, e.g. file printing, it is virtually impossible to ensure Exactly-once Semantics

Server strategy One of two:

1. Send an **ACK** after printing
2. Send an **ACK** before printing

Client strategy One of four:

1. Never resend the request
2. Always resend the request
3. Resend the request when it receives an **ACK**
4. Resend the request when it does not receive an **ACK**

Server Faults and Exactly-once Semantics

Scenario Server crashes and quickly recovers so that it is able to handle client retransmission, but **it has lost all state**

Let

A: ACK

P: print

C: crash

Fault scenarios (ACK \rightarrow P)

1. A \rightarrow P \rightarrow C
2. A \rightarrow C (\rightarrow P)
3. C (\rightarrow A \rightarrow P)

Fault scenarios (P \rightarrow ACK)

1. P \rightarrow A \rightarrow C
2. P \rightarrow C (\rightarrow A)
3. C (\rightarrow P \rightarrow A)

Client

Reissue Strategy

Always
Never
When Ack
When not Ack

OK = Text printed once

Strategy A \rightarrow P

APC AC(P) C(AP)

Dup	OK	OK
OK	Zero	Zero
Dup	OK	Zero
OK	Zero	OK

Dup = Text printed twice

Server

Strategy P \rightarrow A

PAC PC(A) C(PA)

Dup	Dup	OK
OK	OK	Zero
Dup	OK	Zero
OK	Dup	OK

Zero = Text not printed at all

Conclusion No combined strategy works on every fault scenario


At-least-once vs. At-most-once

- Consider a locking service using two RPCs:

`lock()`

`unlock()`


No failures and no message loss



- It does not matter the semantics supported by the RPC library


At-least-once vs. At-most-once: Lost Response

At-least-once




- ▶ Remote procedure may be invoked more than once
 - ▶ If procedure is not **idempotent**:
 - ▶ RPC must include an id as argument
 - ▶ Server must keep table with responses

At-most-once



- ▶ The RPC middleware itself ensures that the procedure is not executed more than once


At-least-once vs. At-most-once: Client crash



- Again, the RPC semantics is irrelevant

At-least-once vs. At-most-once: Server crash


At-most-once



- ▶ Client does not know if server granted it the lock
 - ▶ Depends on when the server crashed
- ▶ Client, **not RPC**, may ask the server (or just retry)
 - ▶ Server needs to remember state across reboots
 - ▶ E.g. store locks state on disk

At-least-once vs. At-most-once: Server crash

At-least-once



- ▶ Server may run the procedure several times
 - ▶ Client stub may send several requests before giving up
- ▶ Server needs to remember previous requests across reboots (if requests are not **idempotent**). E.g.:
 - ▶ Store table request ids on disk
 - ▶ Check the request table on each request

At-least-once vs. At-most-once: Conclusions

Message loss

At-least-once

- ▶ Suits if requests are idempotent

At-most-once

- ▶ Appropriate when requests are not idempotent

Server crashes

- ▶ No clear advantage: the service itself may have to take special measures

Roadmap

Idea

Implementation

Transparency

RPC Semantics in the Presence of Faults

Further Reading

- ▶ Tanenbaum e van Steen, *Distributed Systems, 2nd Ed.*
 - ▶ Section 4.2 *Remote Procedure Call*, except subsection 4.2.4
 - ▶ Subsection 8.3.2 *RPC Semantics in the Presence of Failures*
- ▶ Birrel and Nelson, "*Implementing Remote Procedure Calls*", ACM Transactions on Computer Systems, Vol. 2, No. 1, February 1984, Pages 39-59
- ▶ A. Spector, "*Performing Remote Operations Efficiently on a Local Computer Network*", Communications of the ACM, Vol. 25, No. 4, April 1982, Pages 246-260