

# Clients and Servers (Processing)

March 7, 2018

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

# Roadmap

## Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures


Security

Communication Channel Adaptation


Further Reading

# Clientes e Servidores

- ▶ Most distributed applications have a **client-server** architecture:



- ▶ We'll use *client* and *server* in a broad sense:



- ▶ A server can also play the role of client of another service.

# Roadmap

Definition

**Server/Object Location**

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

# Server/Object Location

**Problem:** how does a client find a server?

**Solution:** not one, but several alternatives:

- ▶ hard coded, rarely;
- ▶ program arguments: more flexible, but ...
- ▶ configuration file
- ▶ via *broadcast/multicast*;
- ▶ via location/naming server (later in the course)
  - ▶ local, like `portmapper` or `rmiregistry`;
  - ▶ global.

# Roadmap

Definition

Server/Object Location

**Distribution Transparency**

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading


# Distribution Transparency

**Issue:** Many distribution transparency facets can be achieved through client side **stubs** (also called **clerks**):

**Access** e.g. via RPC;

**Location** e.g. via multicast;

**Replication** e.g. by invoking operations on several replicas:



**Faults** e.g. by masking server and communication faults

► if possible



# Roadmap

Definition

Server/Object Location

Distribution Transparency

**Concurrency**

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

# Concurrency

- ▶ There are several reasons for using concurrency:
  - ▶ Performance (+ on servers);
  - ▶ Usability (+ on clients) – still performance, really.
- ▶ The goal is to overlap I/O with processing
- ▶ Example: Web service

## Client-side

- ▶ A Web page may be composed of several *objects*
- ▶ A browser can render some objects, while it fetches others via the net.

## Server-side

- ▶ May serve several requests simultaneously



src:Pai et al. 99

# How to Achieve Concurrency?

## Threads

- ▶ Remember SO ...

## Events

- ▶ Remember LCOM ...


# Iterative Web Server



src:Pai et al. 99


- ▶ Has only one thread
- ▶ Processes a request/connection at a time
- ▶ To read the file the server may have to go to disk. In that case, it:
  - ▶ will block and
  - ▶ cannot process other requests
- ▶ Such a server can process only a few requests per time unit

# Multi-threaded Server



src:Pai et al. 99

- ▶ Each thread processes a request (and HTTP 1.0 connection)
- ▶ When one thread blocks on I/O
  - ▶ Another thread may be scheduled to run in its place.
- ▶ A common pattern is:
  - One **dispatcher** thread, which accepts a connection request
  - Several **worker** threads, each of which processes all the requests sent in the scope of a single connection



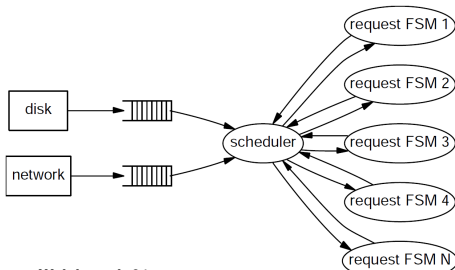
src:Welsh et al. 01

# Event-driven Server



src:Pai et al. 99

- ▶ The server executes a loop, in which it:
  - ▶ waits for events (usually I/O events)
  - ▶ processes these events (sequentially, but may be not in order)
- ▶ Blocking is avoided by using **non-blocking** I/O operations
- ▶ Need to keep a FSM for each request
  - ▶ The loop dispatches the event to the appropriate FSM
- ▶ Known as the state machine approach



src:Welsh et al. 01

# Thread vs. Event Debate

Ease of programming  
Performance

# Thread-based Concurrency: Ease of Programming

- ▶ Appears simple:
  - ▶ Structure of each thread similar to that of an iterative server
  - ▶ Need **only** to ensure **isolation** in the access to shared data structures
- ▶ Could use only monitors and condition variables, e.g. synchronized methods in Java
  - ▶ Not so easy: there are some implications in terms of modularity ([Ousterhout96](#))
  - ▶ Possibility of deadlocks
- ▶ Performance may suffer
  - ▶ The larger the critical sections, less concurrency
  - ▶ But the main reason for concurrency is performance




# Event-based Concurrency: Ease of Programming

- ▶ Programmer needs to:
  - ▶ Break processing according to potentially blocking calls
  - ▶ Manage the state explicitly (using state machines), rather than relying on the stack
- ▶ The structure of the code is very different from that of the iterative server
- ▶ No nasty errors like race conditions, which may be elusive
- ▶ But many complain about lack of support by debugging tools
- ▶ ... and **others** that the it leads to poorly structured code
  - ▶ The author points out that the issue is preemption rather than multithreading
  - ▶ Actually, the problem is **lack of atomicity**
    - ▶ With multiple cores, we can have race conditions, even if there is no preemption

# Thread-Based Concurrency: Performance

- ▶ Same file 8 KB reads (no disk accesses)
- ▶ No thread creation
- ▶ "4-way 500MHz Pentium III with 2 GB memory under Linux 2.2.14"




src: Welsh et al. 01

- ▶ As the number of threads increases, the system throughput increases, then levels-off and finally dives
- ▶ Clearly each thread requires some resources
- ▶ There are also issues concerning context switching
  - ▶ Actually, depends on whether user-level or kernel-level threads

# Event-Based Concurrency: Performance

- ▶ Requires non-blocking (also called asynchronous) I/O operations
  - ▶ Otherwise, must resort to multiple threads to emulate asynchronous I/O
- ▶ Allows user level scheduling
  - ▶ The dispatcher may choose which event to handle next

- ▶ Same file 8 KB reads (no disk accesses)
- ▶ Only one thread
- ▶ As the number of requests in a queue increases throughput increases until it reaches a plateau



src: Welsh et al. 01

- ▶ Needs multiple threads to achieve **parallelism** in multi core/processor platforms

# TB vs EB Concurrency: Performance

- ▶ The debate was somewhat "muddled" by implementations that were less than optimal
- ▶ Actually, at the technical level this is very similar to the debate about user-level vs. kernel-level threads
- ▶ User-level threads are more efficient than kernel-level threads
  - ▶ Function calls vs. system calls
  - ▶ But efficient implementations require OS support for non-blocking I/O
- ▶ But there are some unavoidable blocking, e.g. page faults
- ▶ We need kernel-level threads in order to take advantage of multiple processors/cores

# Server Architectures

Architecture	Paral.	I/O Oper.	Progr.
Iterative	No	Blocking	easy
Multi-threaded	Yes	Blocking	races
State-machine	Yes	Non-blocking	event-driven

- ▶ To take advantage of multiple processors/cores we need to use *kernel-level threads* (or processes).
  - ▶ On state-machine designs we may use multiple threads

# TB vs EB Concurrency: Conclusion

- ▶ Pure thread-based and event-based designs are the extremes in a design space
- ▶ Threads are not as heavy as processes, but they still require resources
  - ▶ You may want to bound their number
- ▶ If you want more parallelism, you need to use both:  
**Threads** virtually all processors now-a-days are multicore;  
**Events** to limit the number of threads, and therefore their overhead
- ▶ There are many frameworks supporting event-driven designs
  - ▶ Java itself offers Java NIO (non-blocking I/O)
  - ▶ Not sure about their performance
    - ▶ They are often built on top of a stack of multiple layers

# Thread-based Concurrency: Basic Considerations

## Java

- ▶ Assume that the Java socket API is not thread-safe
  - ▶ The documentation is mute about this
  - ▶ Java runs on top of different OS
- ▶ You must handle concurrency explicitly

## POSIX

- ▶ It requires many system calls, such as `accept`, `read/write`, `sendto/receivefrom`, to be **thread-safe**
  - ▶ But, data of concurrent `write`'s may be interleaved
    - ▶ I.e., `write/read` may not be **atomic** (apparently it depends on the buffer size)
- ▶ What about `send(to)/receive(from)` ?
  - ▶ When used on `STREAM` sockets, may behave similarly to `write`
  - ▶ When used on `DATAGRAM` sockets, one expects POSIX-atomicity to be implied, but ...
- ▶ To be on the safe side, handle concurrency explicitly

# Thread-based Concurrency: Java

Thread class/Runnable interface for creating threads

- ▶ You can use also thread pools via the interfaces

`java.util.concurrent.ExecutorService` and/or

`java.util.concurrent.ScheduledExecutorService`

Synchronized methods allow for coarse grained CC, similar to monitors

`java.util.concurrent.locks` package for synchronization objects (locks and condition variables) to prevent race conds

- ▶ Check also the `java.util.concurrent.Semaphore`
- ▶ Some classes of the `java.util.concurrent` such as `ConcurrentHashMap` provide a thread-safe version of corresponding `java.util` collection classes

Oracle's Java Tutorials' Concurrency Lesson Overview of core classes

- ▶ For a more practical oriented tutorial you can checkout [java.util.concurrent - Java Concurrency Utilities](#)



# Event-based Concurrency with `java.nio` package

## Core classes

**Channels** There are several subclasses

**Selector** For blocking waiting for more than one I/O event from a selectable channel

**Buffers** To read/write data from/to channels

**Issue** `java.nio.channels.FileChannel` is not selectable

- ▶ To avoid blockin on file I/O need to use `java.nio.channels.AsynchronousFileChannel`, which supports asynchronous I/O
  - ▶ This is more complicated than non-blocking I/O
  - ▶ There is no `java.nio.channels.AsynchronousDatagramChannel`, although one can find references to it on the Web

## Getting started with new I/O (NIO) Overview of Java I/O

- ▶ Refers to non-blocking I/O as asynchronous I/O, but they are not the same
- ▶ For (an even) more practical oriented tutorial you can checkout [Java NIO Tutorial](#)


# Event-driven Server Design by Doug Lea

## Doug Lea's design

- ▶ This is a presentation :(
- ▶ To fill in the details check the [Architecture of a Highly Scalable NIO-Based Server Blog](#)

# Server Clusters


- ▶ In order to support Internet-wide services, we need to use **server clusters/server farms**.
- ▶ A simple approach is to route the requests at the TCP level:



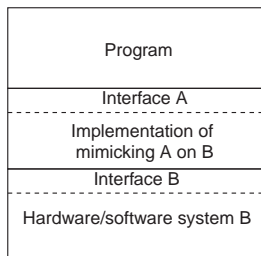
- ▶ The crux is to balance the load on the different servers
  - ▶ **Round-robin** is perhaps the simplest approach
  - ▶ Application-layer solutions are also possible

# Resource Virtualization (1/3)

## Idea



(a)




(b)

Essentially, virtualization allows to emulate the behavior of a different system.


**Objective** Allow legacy SW developed for a specific mainframe architecture to execute on a different platform (IBM)

# Resource Virtualization (2/3)

## Implementations



(a)




(b)

**Process Virtual Machine (a)** Each VM provides an instruction set and a run-time to support the execution of a single application. Example: *Java Virtual Machine*.

**Virtual Machine Monitor (b)** Each VM provides an instruction set and a run-time that allows the concurrent and independent execution of multiples OSs. Example: VMware, Xen, VirtualBox.

# Resource Virtualization (3/3)



- ▶ The use of VMs in distributed systems has mainly two advantages:
  1. Improves security and reliability.
  2. Simplifies management

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

**Keeping State in Servers**

Failures

Security

Communication Channel Adaptation

Further Reading

# Servers and State

**Problem** the execution of the same task on every request may unnecessarily tax the server

**Solution** the server can keep some **state**, i.e. information about the status of ongoing interactions with clients;

- ▶ the size
- ▶ the processing demands

of each message are potentially smaller

- ▶ For example, in a distributed file system, the server may avoid open and close a file for each remote read/write operation
  - ▶ The server may keep a cache of open files for each client
- ▶ Depending on whether or not a server keeps state information, a server is called **stateful** or **stateless**, respectively



# Stateless File Server

- ▶ Consider a simple file service that supports two operations:
  - ▶ read data (from file)
  - ▶ write data (to file)
- ▶ If the server is stateless it keeps no information, therefore each request must include at least:
  - ▶ operation
  - ▶ client id
  - ▶ full path name
  - ▶ position on the file
  - ▶ number of bytes to transfer
  - ▶ data (only in write requests)

Upon a read request the server must:

1. Check permissions for client
2. Open the file (`open()`)
3. Move the file offset to the requested position (`lseek()`)
4. Read the data from the file (`read()`)
5. Close the file (`close()`)

# Stateful File Server

- ▶ Server may keep information on a table about previous requests of each client (e.g.):
  - ▶ file name (or file descriptor)
  - ▶ client permissions
  - ▶ current position (or none)
  - ▶ id of previous request
- ▶ Server may support two additional operations:
  - ▶ open file, which returns a **file handle**
  - ▶ close file
- ▶ Read/write requests need to include only:
  - ▶ operation
  - ▶ operation id
  - ▶ file handle
  - ▶ number of bytes to transfer
  - ▶ data (only in write requests)

Upon a read request the server must:

1. Look up the file handle on the table, to get the file descriptor
2. Read the data from the file (`read()`)

# Stateful Servers and Failures

- ▶ Keeping state information raises some challenges:

- ▶ of consistency;
- ▶ of resource management;

**upon failure** of either clients or server

- ▶ Loss of state when a server crashes may lead to:
  - ▶ ignoring or rejecting client requests after recovery:
    - ▶ the client will have to start a new **session**
  - ▶ wrong interpretation of client requests sent before the crash:
    - ▶ TCP connection port reuse
- ▶ Keeping state (on server) when the client crashes may lead to:
  - ▶ resource depletion
    - ▶ E.g. if a client crashes before invoking `close()`
  - ▶ wrong interpretation of requests sent by other clients after the crash
    - ▶ If client id is reused (e.g. IP address and port number)

# Stateful Servers and Client Crashes

**Challenge** resources reserved for the client may remain allocated forever

- ▶ sockets, for connection based communication
- ▶ state, in the case of stateful servers
- ▶ application specific resources

**Solution** **leases** (and timers):

- ▶ a server *leases a resource* to a client for only during a finite time interval: upon its expiration, the resource may be taken away, unless the client **renews** the lease

# Stateless Servers and Message Loss

- ▶ Stateless servers are not immune to problems arising from failures:
  - ▶ message duplication may lead to handling the same request several times
    - ▶ operations must be **idempotent**, if the transport protocol does not ensure non-duplication of packets;
    - ▶ even if the transport protocol ensures non-duplication of packets, we may still need idempotent operations
      - What if the connection breaks?
- ▶ How can stateful servers handle duplicated requests?
  - ▶ Need to be careful about client identification

# Stateful Servers and Client Identification

1. Use the address of the **access point**, i.e. of the channel endpoint
  - ▶ For example, the client's IP address and port
  - ▶ Issue: may not be valid for more than one transport session:
    - ▶ E.g. if a TCP connection breaks and a new one is setup in its place, the port number on the client's side may be different
2. Use a transport-layer independent **handle**. For example:
  - ▶ HTTP cookies

# Servers, State and Protocols

- ▶ **Obs.-** Statelessness is a protocol issue:
  - ▶ A server can be stateless only if each protocol message has all the information for its processing independently of previous communication;
  - ▶ Likewise, a server can be stateful only if each protocol message has enough information to relate it to previous communication
- ▶ For example, Netscape had to add HTTP-header fields specifically for cookies.
  - ▶ HTTP is essentially stateless
    - ▶ Version 1.0 even used one TCP connection per request
  - ▶ Cookies are a device that allows a server to keep state about a client session (and can also be abused):
    - ▶ *cookies* are stored on the client side

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

**Failures**

Security

Communication Channel Adaptation

Further Reading



# Failures

## Challenges:

1. components in a distributed application may fail, while others continue operating normally
2. on the Internet it is virtually impossible to distinguish network failures from host failures or even a slow host

**Solution:** highly application dependent, but we'll study some general techniques

## Distribution is harder than concurrency

In concurrent (local) systems the programmer needs to consider all possible execution interleavings

In distributed systems the programmer needs **also** to consider all possible failures

- Distributed systems are inherently concurrent

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

**Security**

Communication Channel Adaptation

Further Reading

# Security

**Challenge:** servers execute with privileges that their clients usually do not have

**Solution:** servers must

**authenticate clients:** i.e. "ensure" that a client is who it claims to be;

**control access to resources:** i.e. "ensure" that a client has the necessary permissions to execute the operation it requests.

- ▶ A related requirement is data **confidentiality**
  - ▶ need to encrypt data transmitted over the network
- ▶ Code migration (i.e. downloaded from the network) raises even more issues.

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

**Communication Channel Adaptation**

Further Reading

# Communication Channel Adaptation

**Order** the application will have to reorder the messages (must use a sequence number), if that is important

**Reliability** need to use timers to recover from message loss. Have to be aware of the possibility of duplicates.

**Flow control:** if you want to avoid message loss because of insufficient resources

**Channel abstraction:** the application may have to build messages from a stream. Or, fragment messages at one end and reassemble them at the other end.

# Roadmap

Definition

Server/Object Location

Distribution Transparency

Concurrency

Keeping State in Servers

Failures

Security

Communication Channel Adaptation

Further Reading

# Further Reading

- ▶ Ch. 3 of Tanenbaum e van Steen, *Distributed Systems, 2nd Ed.*
  - ▶ Subsection 3.1.2 *Threads in Distributed Systems*, we assume the remaining material in Section 3.1 to be background knowledge (OS class)
  - ▶ Subsection 3.3.2 *Client-Side Software for Distribution Transparency*
  - ▶ Section 3.4 *Servers*
  - ▶ Section 3.2 *Virtualization*
- ▶ Arpaci-Dusseau & Arpaci-Dusseau, *Event-based Concurrency*, Ch. 33 of OSTEP book
- ▶ Pai et al., *Flash: An efficient and portable Web Server*, in 1999 Annual Usenix Technical Conference
- ▶ Welsh et al, *SEDA: An Architecture for Well-Conditioned, Scalable Internet Services*, in Symposium on Operating Systems, 2001