

# Names

March 21, 2018

# Outline

The Problem

Naming Concepts

Name Spaces

Name Resolution

Additional Reading

# Server/Object Location

**Problem:** How does a client know where is the server?

**Solution:** Not one, but several alternatives:

- ▶ *hard coded*, seldom;
- ▶ via program arguments: more flexible, but ...;
- ▶ via configuration file;
- ▶ via *broadcast/multicast*;
- ▶ via a location/name service:
  - ▶ local, e.g. *rmiregistry*.
  - ▶ *global*.

# Addresses

- ▶ In the labs, we have used IP addresses (and ports)
- ▶ Addresses are **names** of **access points**
- ▶ Addresses have some limitations:
  - ▶ Addresses often are location dependent and change frequently
    - ▶ E.g. when a service is moved from one computer to another
  - ▶ Addresses are bound to access points, making it difficult, e.g., to balance the load among several servers.

# Names

- ▶ Names are ... sequences of symbols (bits/characters/...) that refer to entities/objects. Or better (Schoch):

*The **name** of a resource indicates **what** we seek,  
an **address** indicates **where** it is,  
(and a **route** tells us **how** to get there.)*

- ▶ Names have some advantages over addresses:
  - ▶ They can be human-friendly.
  - ▶ They can hide both complexity and dynamics
    - ▶ E.g. they can hide access point changes
- ▶ Naming is a layer of indirection
  - ▶ Ultimately you need an address to access/operate on an object
  - ▶ But, as someone said:  
*"All problems in CS can be solved by another layer of indirection"*

# Pure Names

- ▶ Are names that contain no information whatsoever about what they refer to:
  - ▶ Not only about location, but about anything else
  - ▶ They do not commit the system to anything
  - ▶ They are useful only for comparison

## Problems/challenges of pure names

- ▶ where to look them up to find out information about them?
- ▶ how do you know that an object does not exist? How can a global search be avoided?
- ▶ how to engineer uniqueness reliably in a distributed system?

## Problems/challenges of impure names

- ▶ what if the information yielded by the name, e.g. location, is not valid anymore?
  - ▶ This is specially relevant for mobile systems, and requires appropriate solutions
- ▶ Examples?
  - ▶ From the "real" world?
  - ▶ From the "virtual" world?

# Identifiers

- ▶ An **identifier** is a name with 3 properties:
  1. an identifier refers to an entity at most;
  2. an entity has at most one identifier;
  3. an identifier refers always to the same entity (it is never reused).
- ▶ Identifiers provide a mean to refer to an entity in a precise way, independently of its access points.
- ▶ Examples?
  - ▶ From the "real" world?
  - ▶ From the "virtual" world?

# Name Spaces

## Name Space

- ▶ Names are organized into name spaces.
- ▶ A name space is composed of a set of bindings of a name to an object/entity (or to one of its attributes, e.g. addresses)
- ▶ A name space defines:
  - ▶ the syntax and structure (flat vs. hierarchical) of a name in that space;
  - ▶ the rules to find a binding of a name (name resolution)

## Naming Domain/Context

- ▶ Is the scope of a name
  - ▶ Names in a domain/context must be unique
- ▶ Allows to partition a name space
  - ▶ Often, a naming domain has an administrative authority for assigning names within it
  - ▶ An administrative authority may delegate name assignments for sub-domains (e.g. in DNS)



# Name Space Structure

- ▶ Most name spaces have a hierarchical structure:
  - ▶ OS filesystem
  - ▶ Domain Name System (DNS)
  - ▶ Postal addresses
  - ▶ Car license plates are resolved in another context – per country, region etc.
- ▶ A hierarchical structure simplifies:
  - ▶ the assignment;
  - ▶ the resolutionof names
- ▶ Up to recently (about the year 2000) nobody knew how to implement efficiently name resolution in a flat name space.


# Name Resolution - binding

- ▶ **Name resolution/binding** is the procedure to find the value for an attribute, usually an address, of a named entity that allows that entity to be used/accessed
- ▶ A name is always resolved in the context of its name space:
  - file name -> OS filesystem
  - Java program variable -> JVM executing the program
  - ISBN of a publication -> ISBN (Intern.Standard Book Number)
  - Car license plates -> national/regional license plate register
- ▶ Late binding is considered good practice
  - ▶ Programs should use names rather than addresses
  - ▶ Name resolution is performed by a naming service (e.g. `rmiregistry`)


# Name Resolution in a Distributed System

- ▶ In small scale distributed systems, name resolution requires only one server:
  - ▶ E.g., the `rmiregistry` or the `portmap`.
- ▶ In distributed systems of larger scale, name resolution may require more than one server. In this case, name resolution can use one of 3 strategies:
  1. Iterative
  2. Recursive.
  3. Transitive.


# Name Resolution: Strategies



**Iterative**



**Recursive**



**Transitive**

- ▶ Recursive name resolution:
  - ▶ Allows for caching at servers
    - ▶ This may make resolution more efficient (with lower communication costs)
  - ▶ But, it:
    - ▶ requires servers to keep state
    - ▶ makes it harder to set the values of timeouts (unless there is immediate acknowledgment)
- ▶ Transitive name resolution is seldom used

# Name Resolution and *Closure Mechanism*

Names are resolved always in a context

## Problem

- ▶ How do you get an initial "context" from which you can start resolving a names?
  - ▶ How do you get a "remote reference to the `rmiregistry`"?
  - ▶ How to start the name resolution of a name of a file system:  
i.e. where is the root directory?
  - ▶ How to find the IP address of a DNS server to resolve a DNS name?

## Response

Use a **closure mechanism**

- ▶ Typically this is an *ad-hoc* and simple solution.

# Additional Reading

- ▶ Chapter 5 of Tanenbaum e van Steen, *Distributed Systems, 2nd Ed.*
  - ▶ Section 5.1: *Names, Identifiers and Addresses*
  - ▶ Section 5.3: *Structured Naming*