

# Name Resolution in Flat Name Spaces

## Distributed Hash Tables (DHTs)

Pedro F. Souto (`pfs@fe.up.pt`)

April 3, 2018

# Resolution of Unstructured Names

## Problem

- ▶ Assume you want to develop a "peer-to-peer" version of the backup service on the Internet.
- ▶ How do you locate the peers storing a given chunk of a file?
  - ▶ Each file has a 256-bit id
  - ▶ This id is **unstructured**

## No solution Broadcasting/multicasting

- ▶ It just does not scale beyond a LAN

## Issue How do we **resolve** efficiently an unstructured name on the Internet?

## Solution Use a distributed hash table (DHT)

- ▶ Answer provided by academia to the problem of locating an entity in P2P system

# Distributed Hash Table (DHT)

- ▶ A DHT is similar to a **hash-table**
  - ▶ It maps a **key** to **value**
  - ▶ The **key** is an object identifier
  - ▶ The **value** is an address
    - ▶ assume it is the address of the node/peer **responsible** for the key
- ▶ A DHT provides a single operation:  
`lookup(key)` returns the address of the node responsible for the key
  - ▶ The address can be used to insert an object, to access to an object ...
- ▶ In a DHT-based system, node identifiers and key values are drawn from the same domain, e.g. a number with  $m$  bits
- ▶ The node responsible for a key value is the one whose identifier is **closer** to that key
  - ▶ Depending on the definition of **distance** we get different DHTs


# DHT Example: Chord

- ▶ Chord uses identifiers with  $m$ -bits ordered in a ring ( $\text{mod } 2^m$ )
- ▶ Each "object" has an  $m$ -bit random identifier: the key of DHT entries ( $m = 128$  in the original paper - used MD5)
- ▶ Each node has an  $m$ -bit random identifier
  - ▶ This is different from a key value, which is an address

- ▶ The node **responsible** for key  $k$  is the **successor** of key  $k$ ,  $\text{succ}(k)$ :

$\text{succ}(k)$  is the node with the **smallest** id that is larger or equal to  $k$  ( $\text{succ}(k) \geq k$ )

- ▶ Given a key  $k$  the node responsible for it will have an id **higher or equal** to  $k$ .



src: Stoica et. al. 2001

# Key Resolution in Chord (1/2)

**Problem** Given a key  $k$ , how do you find  $\text{succ}(k)$ ?

**No Solution 1** Each node  $n$  keeps information about the next node in the ring ( $\text{succ}(n + 1)$ )


- ▶ Can use any resolution strategy (iterative, transitive or recursive)
- ▶ ... but it does not scale. Why?

**No Solution 2** Each node  $n$  keeps information about all nodes in the ring

- ▶ Constant time name resolution
- ▶ ... but it does not scale. Why?

## Key Resolution in Chord (2/2)

**Solution** In addition to a pointer to the next node in the ring each node keeps pointers that allow it to reduce at least in half the **distance** to the key



► Because nodes that are  $2^i$  apart may not be active, each node  $n$  keeps a pointer to the  $\text{succ}(n + 2^i)$  for  $i = 0 \dots m - 1$

- This scheme has 3 important properties:
1. Each node keeps information on only  $m$  nodes
  2. Each node knows more about nodes closer to it than nodes farer away
  3. The table in a node may not have information on the  $\text{succ}(k)$ , for some  $k$  – i.e. a node may be unable to resolve a key by itself
  4. But key resolution requires  $O(\log(n))$  steps


## Chord: *Finger Table* (1/2)

- ▶ The ***Finger table***,  $FT_n[]$ , is an array with  $m$  pointers:

$$FT_n[i] = \text{succ}(n + 2^{i-1}) \bmod 2^m \text{ where } i = 1 \dots m$$

- ▶  $FT_n[1]$  is the node that follows  $n$  in the ring
- ▶ To resolve (*lookup*) a key  $k$ , node  $n$  forwards the request to:
  - ▶ The next node, i.e.  $FT_n[1]$ , if  $n < k < FT_n[1]$
  - ▶ To node  $n'$  st  $n' = FT_n[j] \leq k < FT_n[j + 1]$   
(All arithmetic in modulo  $2^m$ )
- ▶ Chord works correctly iff  $FT_n[1]$  is correct
  - ▶ Chord tolerates transient inconsistencies in other elements of  $FT_n[]$ , by trying the resolution again (may not be necessary even)
- ▶ The original Chord paper describes an iterative resolution scheme
  - ▶ Allows to update the *Finger Table*.

# Chord: *Finger Table* (2/2)





# Chord: Other Issues

**Node Joining** Node  $n$  can ask any node to locate  $\text{succ}(n)$

- ▶ The crux is to get the  $FT_x[1]$  correct
- ▶ This process can be simplified if each node keeps a pointer to its predecessor
- ▶ Periodically, a node sends a message to the next node in the ring and updates its finger table

**Node Failure** Rather than keep a single successor, a node keeps a list of  $r$  successors

- ▶ If the successor fails, a node can replace it with another one from that list

**Identifiers Generation** To achieve some tolerance to denial-of-service (DoS) attacks, identifiers should be generated using a cryptographic hash function, e.g. SHA256

# Virtual Topology Issues (1/2)

**Problem** Chord, and other P2P systems, use an overlay network

- ▶ If the topology of the overlay network is oblivious to the underlying physical network, routing of messages along the overlay network may be inefficient
  - ▶ Messages may follow an erratic route, e.g. bouncing between hosts in different continents

**Sol. 1: Assign identifiers according to the underlying topology**

- ▶ I.e. assign identifiers so that the overlay topology is close to that of the underlying physical topology.
- ▶ This is not always possible. E.g. it is **not** possible in Chord.

# Virtual Topology Issues (2/2)

## Sol. 2: Route messages according to the underlying topology

- ▶ For example, Chord could keep several nodes per interval  $[n + 2^{i-1}, n + 2^i]$  rather than a single one, and when resolving a key, might use the closest node

## Sol. 3: Pick neighbors according to the underlying topology

- ▶ In some algorithms, nodes can pick their neighbors, i.e. establish the links of the overlay network.
- ▶ This is not always possible. E.g. it is **not** possible in Chord.

# Further Reading

- ▶ Subsection 5.2.3, Tanenbaum and van Steen, *Distributed Systems*, 2nd Ed.
- ▶ I. Stoica et al., "Chord: A scalable peer-to-peer lookup protocol for Internet applications", *IEEE/ACM Transactions on Networks*, (11)1:17-32, Feb 2003 (acessível via biblioteca digital da ACM “dentro da FEUP”)